

# Assignment 1: Neural Network Foundations with JAX

**Due:** February 13, 2026

**Submission:** Brightspace Assignment 1

## Objective

Understand optimization and architecture design using JAX. Learn how JAX's functional programming paradigm differs from PyTorch/TensorFlow and explore the role of `jit` and `grad` in performance and automatic differentiation.

## Tasks

1. **Implement a simple feedforward neural network from scratch using JAX primitives**
  - No high-level libraries like Flax or Haiku for the core logic.
  - Include parameter initialization, forward pass, loss computation, and gradient-based updates using `jax.grad`.
2. **Train your network on the MNIST dataset**
  - Use an appropriate training/validation/testing split.
  - Normalize inputs and one-hot encode labels.
3. **Compare performance with a PyTorch or TensorFlow implementation**
  - Implement the same architecture in PyTorch or TensorFlow.
  - Measure:
    - First epoch time (including JAX compilation overhead)
    - Steady-state epoch time
    - Final test accuracy

- **Batch size comparison:** Train with batch sizes of 64, 256, and 1024. Record timing and accuracy for each.
- 4. **Explain in a short report (maximum 2 pages): Use results from your experiments to support the discussion below**
  - How does JAX’s functional programming paradigm differ from PyTorch/TensorFlow?
  - Discuss the role of `jit` and `grad` in performance and automatic differentiation.
  - **Focus on why performance differs rather than expecting big speedups:**
    - Compilation overhead in JAX (`jit` warm-up cost)
    - Functional style and immutability
    - Graph optimizations and kernel fusion
    - Impact of batch size on amortizing compilation cost
  - Reflect on your experience: What was easier/harder in JAX vs PyTorch/TensorFlow?

## GPU Requirement

- **You must run your experiments on a GPU** to ensure realistic performance comparisons.
- Options:
  - **Morningstar cluster** (available through the university)
  - **Your personal GPU**
  - **Google Colab** or other cloud-based GPU services (e.g., Kaggle, AWS, Azure)
- Document which platform you used and any setup steps in your report.

## Deliverables

- **Code:**
  - JAX implementation (pure JAX, no Flax)
  - PyTorch/TensorFlow implementation
  - Reproducibility readme file

- **Report (PDF):**
  - Architecture details
  - Performance comparison (include timing and accuracy for different batch sizes)
  - Conceptual discussion
  - Reflection on AI tool usage (if applicable): *What did AI suggest? What did you accept/reject and why?*

## Grading Rubric (100 points)

Component	Points
<b>JAX Implementation</b>	30
- Correct forward pass & loss	10
- Gradient updates using <code>grad</code>	10
- Use of <code>jit</code> for optimization	10
<b>PyTorch/TensorFlow Implementation</b>	20
- Correct architecture & training	10
- Performance comparison included	10
<b>Report (Conceptual Discussion)</b>	30
- JAX vs PyTorch/TensorFlow paradigm	10
- Role of <code>jit</code> and <code>grad</code> explained	10
<b>Analysis of performance differences</b>	10
<b>Code Quality &amp; Documentation</b>	10
- Clear, readable, commented code	10
<b>Bonus</b>	10
- Visualization of training curves	5
- Insightful AI tool reflection	5